

Case study — Senior Product Manager, AI (Discovery, International)

Why we do it this way

The role is about building with AI, not writing about building with AI. So we ask you to build something that runs, and something we can click through. Not a deck, not a memo with screenshots.

There is no time limit and we mean it. Spend two hours or spend twelve. We ask you to log the time and tell us which AI tools you used, because how much you get done per hour is genuinely part of what we are assessing. More hours is not a better score. A sharp result in three hours beats a sprawling one in twenty, and we would rather see where you chose to stop than see everything you could think of.

If anything here is ambiguous, decide, and tell us what you decided. We will not answer clarifying questions on scope — how you handle an underspecified brief is part of the exercise. Mail us if something is genuinely broken.

The situation

Groupon's international markets are about a third of our orders but under a quarter of our revenue. Average order value outside the US runs at roughly 60% of US levels. Customers are finding us and buying. They are just not finding the right thing often enough.

Our search, browse and homepage were built for the US: dense supply, English-language queries. The same platform runs in Germany, France, Spain, Poland and the UK against thinner inventory and five different languages.

We have given you a month of search logs from five international markets and the deal inventory those searches ran against.

The data

Both files are in this folder. Download them rather than opening them in Sheets — they are plain CSV.

search_log.csv — 8,997 searches, June 2026

Column	Meaning
query_id	Unique ID
market	GB, DE, FR, ES, PL
city	City the customer was searching in
raw_query	What they typed, verbatim and uncleaned
results_shown	How many deals the current system returned
clicked	Did they click any result
purchased	Did that session end in a purchase
date	Search date

deals.csv — 568 live deals across the same markets and cities, with title, category, price, rating and whether the deal is instantly bookable.

The data is synthetic but it is modelled on real patterns and real market weights. Deal titles are in local languages. Queries are as customers typed them, including in languages that do not match the market.

What we want back

Two parts, and they have to connect.

Part A — What is broken

Your read of the data, with the numbers behind it. Whatever form you like: a notebook, a script, a short analysis doc. We will run it or check it, so show your working.

We are deliberately not telling you what to look for.

Part B — What right looks like

A working prototype we can click through. It should demonstrate the discovery experience you think these customers should get, given what you found in Part A.

Use whatever gets you there fastest. Claude Artifacts, v0, Lovable, Bolt, Replit, hand-written React, anything. **We are not assessing whether you wrote the code yourself.** We are

assessing whether you can get from a data problem to a working thing without waiting for an engineer.

One requirement: **your prototype has to handle every kind of query you found in Part A, including the ones you cannot fix.** We will type things into it. What it does when it has no good answer tells us as much as what it does when it has one.

We are not telling you what to build. Choosing the right thing to build against this data is most of the job.

Part C — A writeup. Two pages maximum.

- What you found, with the numbers
- What you decided to build and why that rather than the other options
- What you would do next with an engineering team, and what you would need from the platform teams who own search globally
- How you would measure whether it worked in production, and what would tell you it was not working

And a short log

Hours spent, roughly. Which AI tools you used and what for. One or two sentences on anything the tools got wrong that you had to catch.

How to send it

A git repo, a zip, or links. For the prototype, a hosted link is easiest for us. Include whatever is needed to run Part A. If it needs an API key, say which and do not include yours.

What we are assessing

- Whether you found what is actually going on in this data, rather than the first thing that looks like a finding
- Whether the prototype reflects what you found, rather than being a nice search UI bolted on beside the analysis
- Whether your claims survive us checking them against the data
- Whether you were honest about the limits of what you built
- How much you got done, per hour

What we are not assessing

Production code quality. Test coverage. Visual design. Whether you hand-wrote a single line of it. Ship the rough thing that works.

If you have a disability or circumstance that makes this format difficult, tell your recruiter. We will find another way to assess the same thing.